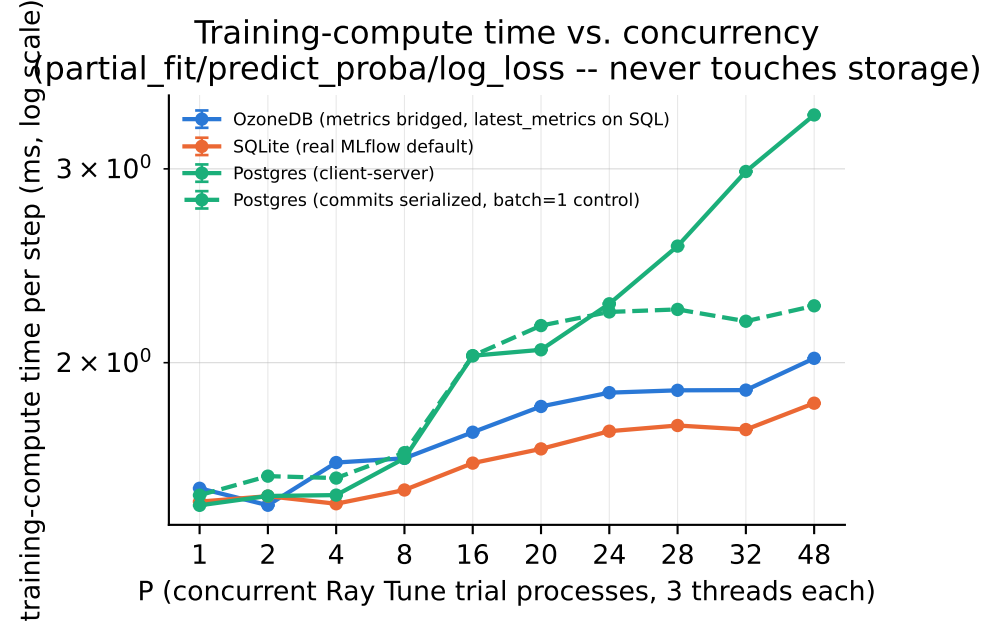
Training-compute time vs. concurrency
(partial_fit/predict_proba/log_loss -- never touches storage)
OzoneDB (metrics bridged, latest_metrics on SQL)
SQLite (real MLflow default)
Postgres (client-server)
Postgres (commits serialized, batch=1 control)
3×10^0
2×10^0
training-compute time per step (ms, log scale)
1
2
4
8
16
20
24
28
32
48
P (concurrent Ray Tune trial processes, 3 threads each)